

i3 TECHNOLOGIES

Formally Instrumented Interconnected & Intelligent Technologies Limited

IBM Silver Business Partner | Red Hat Authorised Training Partner | Anthropic Partner

CONFIDENTIAL - BOARD PROPOSAL

MASTER STRATEGIC & TECHNICAL PROPOSAL

SOVEREIGN AI APPLIANCE FOR AFRICA

IBM Power / LinuxONE hardware · watsonx Orchestrate & Sovereign Core
Open-source foundation models + Granite specialty SLMs, packaged per industry vertical

PREPARED FOR

i3 Technologies Ltd — Board & IBM Alliance

GEOGRAPHIC FOCUS

Kenya -> E. Africa -> MEA

DOCUMENT REFERENCE

i3-SAA-MP-2026-09

PREPARED

September 2026

BOARD BRIEFING NOTE

. Document Purpose & Scope

This master proposal investigates how i3 Technologies can package IBM hardware, IBM's watsonx Orchestrate / Sovereign Core software layer, and a curated stack of open-source and Granite specialty small language models (SLMs) into a productised **Sovereign AI Appliance** — sold and delivered the way classic data-warehouse appliances (Netezza, PureData, Exadata) were once sold: as a pre-integrated, pre-tuned unit of compute, orchestration and vertical intelligence, rather than as a custom integration project rebuilt for every customer.

It builds directly on i3's existing internal workstreams — the IBM Bob Implementation Playbook and ROKS/RHOAI AI Lab architecture, the 16-product Vertical AI Solutions Blueprint (AgriGranite, FinGranite, GovGranite and related Granite-based products), and the Royal Partner Program spanning 11+ African countries and the IBM Research Lab Africa collaboration at CUEA — and extends them into a hardware-plus-software-plus-model appliance line that can be shipped to a customer site, collocated at an i3/partner facility, or deployed inside a customer's own data centre.

Sponsor	i3 Technologies Ltd — Executive Board / IBM Alliance
Prepared as	Companion workstream to the SIT × i3 Digital Work-Centers proposal (AI Services Utility layer)
Primary market	Kenya (launch) -> East Africa (Year 1–2) -> Middle East & Africa / MEA (Year 2–3)
Core partners referenced	IBM (Power, LinuxONE, watsonx, Red Hat), Anthropic, Royal Partner Program network
Status of figures	Indicative planning figures for board discussion; not a binding commercial quotation

Prepared: September 2026.

1. Executive Summary

IBM's own 2026 product direction has converged on exactly the problem i3 is positioned to solve for Africa. At Think 2026, IBM made **Sovereign Core** generally available — an infrastructure layer that embeds governance and compliance controls at runtime so regulated organisations can run AI entirely within their own sovereign boundary — and moved the next generation of **watsonx Orchestrate** into an agentic control plane that can govern agents "from any vendor." In parallel, IBM shipped the **Power S1112**, a compact, single-socket Power11 server purpose-built to run AI inference locally, on-premises, "without shipping data anywhere else," and released **Granite 4.1/4.2** as Apache-2.0, cryptographically signed, open-weight small language models designed to be audited, fine-tuned and run entirely inside a customer's network boundary.

At the same time, Kenya and its neighbours are moving from AI policy aspiration to enforceable infrastructure requirement: Kenya's cloud policy already requires government data to be stored within the country, a National AI Strategy and a draft AI Bill are formalising a risk-based regime with a proposed AI Commissioner, and 44 African countries now have data protection laws with 38 functioning regulators. The market conversation across the continent has shifted from "which cloud AI API should we call" to "where does our AI actually run, and who can prove it."

The recommendation of this proposal is that i3 should not sell AI consulting hours, and should not sell bare IBM hardware. i3 should sell a packaged, pre-integrated, factory-tested product line — the i3 Sovereign AI Appliance (i3-SAA) — that bundles IBM Power/LinuxONE compute, watsonx Orchestrate and Sovereign Core, a curated open-source and Granite SLM model garden, and i3's own vertical-industry accelerators (built on the existing AgriGranite/FinGranite/GovGranite blueprint), delivered under three commercial motions: shipped appliance at the customer site, collocated at an i3/partner facility, or deployed inside the customer's own data centre with i3 as the managed-service operator.

This mirrors the commercial logic that made data-warehouse appliances (Netezza, Teradata, PureData, Exadata) durable, high-margin franchises: a fixed, certified bill of materials; a small number of SKUs sized to workload; a repeatable install; and a managed-service annuity layered on top. Applied to sovereign AI, the appliance model turns i3 into Kenya and East Africa's default answer to "we need AI that never leaves our building, and we need it running in weeks, not a year-long integration project."

The appliance line at a glance

Layer	What it is	Primary source
Hardware	Power S1112 / larger Power11 nodes with on-chip MMA and optional Spyre PCIe AI accelerator; LinuxONE option for extreme-resiliency finance/telco workloads	IBM
Orchestration & governance	watsonx Orchestrate (agentic control plane) + IBM Sovereign Core (runtime policy/compliance enforcement) on Red Hat OpenShift AI	IBM / Red Hat
Foundation models	Open-source LLMs (Llama, Mistral-class) for general reasoning, paired with Apache-2.0 Granite 4.x SLMs (micro/tiny/small/vision/speech/guardian) for cost-efficient, auditable enterprise tasks	Open source + IBM Granite
Vertical intelligence	i3's Granite-based vertical product line — AgriGranite, FinGranite, GovGranite and related fine-tuned specialty models mapped to industry workflows	i3 Technologies
Delivery	Shipped appliance, collocated at i3/partner facility, or deployed in customer data centre — always with i3 managed operations wrapped around it	i3 Technologies

2. Why Now: The Market Has Moved Toward Sovereign, On-Prem AI

Three independent trends have converged in 2026 in a way that did not exist even eighteen months ago, and each directly de-risks and strengthens the appliance thesis.

2.1 IBM's own roadmap has become an appliance roadmap

At Think 2026 (Boston, May 2026), IBM announced its most comprehensive expansion of enterprise AI and hybrid cloud management to date, anchored by two flagship releases: **Sovereign Core** — now generally available — which “transforms digital sovereignty from a policy concept into an operational, runtime-enforceable reality,” letting organisations run AI models and agents entirely within sovereign boundaries without depending on hyperscaler-managed regions; and the next-generation **watsonx Orchestrate**, an agentic control plane in private preview that can deploy and govern agents “from any vendor” under one consistent policy layer. IBM's stated launch partners for Sovereign Core already include AMD, Dell, Intel, MongoDB and Palo Alto Networks — the pattern is an ecosystem of certified, governed building blocks, which is precisely the appliance model.

On the hardware side, IBM shipped the **Power S1112** in July 2026 — a compact, single-socket Power11 system “built for organisations that are extending AI workloads to the edge or on-premises locations outside the data center,” using Power11's on-chip Matrix Math Acceleration for local inference. IBM Redbooks documentation for Power11 explicitly frames the platform, paired with the **Spyre** PCIe AI accelerator and Red Hat AI/vLLM, as delivering “on-premises AI with predictable performance and data-sovereign control” for regulated, data-intensive environments — language that could describe the i3-SAA brief almost verbatim.

On the model side, **Granite 4.x** (4.0 through the April 2026 4.1 release and subsequent 4.2 update) is IBM's bet that certified small models, not ever-larger frontier models, will do most enterprise work — released under Apache 2.0, cryptographically signed, and designed so a customer can download the weights, run them entirely on their own infrastructure, and pay only for compute rather than per-token API access. Granite 4.1 spans 3B/8B/30B language variants plus dedicated vision, speech, embedding and guardian (safety) models — exactly the component set an appliance vendor needs to build vertical products rather than a single general chatbot.

2.2 Africa's regulatory ground has shifted from policy to enforceable infrastructure

By early 2026, 44 African countries had data protection laws in force, with 38 operating functional data-protection authorities; Kenya, Ghana, Nigeria and Algeria now require certain categories of data to be stored or processed within national borders. Kenya has published a cloud policy requiring government data to be stored within its own territory, and a March 2026 tabled Artificial Intelligence Bill proposes an Office of the AI Commissioner with authority to mandate audits of high-risk systems and five-year data retention. Kenya's July 2026 draft Artificial Intelligence and Other Emerging Technologies Policy is explicitly structured around nine pillars including “Sovereignty and Strategic Autonomy,” building on the existing National AI Strategy 2025–2030 and the Data Protection Act, 2019.

This is not abstract policy. Kenyan and regional firms are already building the physical layer: two Kenyan companies launched **Servevah Cloud** at the iXAfrica facility in Nairobi in March 2026, described as the country's first sovereign-hosted AI cloud platform, and the Africa data-centre sector as a whole is projected to attract a further USD 8.76 billion in investment by 2031. The African Union's Continental AI Strategy puts data governance and sovereignty at its centre, noting that most African data — from agricultural satellite imagery to patient health records — is still stored and processed outside the continent it is generated on.

“We cannot build an AI economy on foreign soil — the infrastructure must be here,” Kenya's former Cabinet Secretary for ICT and the Digital Economy, Joe Mucheru, told Panapress in 2026 — a statement that, for i3's purposes, doubles as a product brief.

2.3 Enterprise buyers are stalling on cloud-only AI for exactly the reasons an appliance solves

Globally, the sovereign cloud market reached an estimated USD 80 billion in 2026, growing 35.6% year-on-year, driven by enterprises moving regulated AI workloads away from US hyperscalers — a pattern reinforced by extraterritorial exposure concerns such as the US CLOUD Act and tightening frameworks like the EU AI Act's Article 10 data-governance requirements. IBM's own Think 2026 commentary on UK financial-services buyers captured the pattern directly: a bank's CISO “would not sign off on any cloud-only deployment for processes touching financial data” — the on-premises option was “what made the initial conversation possible.” The same governance anxiety exists, arguably more acutely, among Kenyan banks, insurers, SACCOS, hospitals and county governments, none of whom currently have an easy, pre-integrated way to buy it locally. That gap is the opening for i3.

3. Core Thesis: The Sovereign AI Appliance Model

Historically, the data-appliance category (Netezza, Teradata, IBM PureData, Oracle Exadata) won because it converted a hard, bespoke integration problem — tuning a database engine to specific hardware for acceptable performance — into a fixed, factory-built, pre-tested product that a customer could order by SKU and have running in days rather than months. Buyers paid a premium for certainty: known performance, known support boundary, a single vendor accountable end-to-end.

Sovereign enterprise AI in Africa today has the same shape of problem. A bank, insurer, hospital group or county government that wants AI it fully controls currently faces a bespoke integration project: choose a server vendor, choose an orchestration layer, choose and fine-tune models, build the governance and evaluation tooling, and staff a team to run all of it — typically a 9–18 month effort with meaningful execution risk. **The i3 Sovereign AI Appliance collapses that into a catalogue purchase:** a small number of pre-sized, pre-integrated, pre-evaluated units, each pre-loaded with the vertical models a given industry actually needs, shippable, collocatable or deployable in the customer's own facility, with i3 as the single accountable vendor for hardware, software, models and ongoing operations.

The i3-SAA flywheel

Sell certified appliances -> install and operate them -> harvest lawful, governed operational data to improve the vertical models -> productise the improvements as new specialty SLMs -> expand the appliance catalogue -> grow the managed-services annuity -> replicate across counties and East African markets -> export the proven model into MEA. This is the same flywheel logic i3 has already articulated for its AI Services Utility positioning, now anchored to a physical, ownable product rather than a purely services engagement.

Why “appliance,” specifically, and not just “managed cloud AI”

Dimension	Cloud-only AI (status quo)	i3 Sovereign AI Appliance
Data residency	Data typically leaves the customer's facility or country	Data never leaves the customer's building unless they choose to allow it
Regulatory posture	Depends on foreign provider's terms and jurisdiction	Built to Kenya Data Protection Act / ODPC and the emerging AI Bill from day one
Cost model	Per-token, foreign-currency, exposed to FX and price changes	Capex/opex hybrid in local currency, predictable at scale
Latency & availability	Dependent on international connectivity	Local inference; resilient to connectivity outages
Vendor accountability	Multiple vendors, unclear single point of ownership	Single accountable vendor: i3, backed by IBM/Red Hat/Anthropic
Time to production	Bespoke integration, 9–18 months typical	Pre-integrated SKU, weeks to first production workload

4. Reference Architecture

The appliance architecture extends the ROKS/RHOAI AI Lab foundation i3 has already built internally (Terraform-provisioned, Red Hat entitlement-optimised) into a productised stack with five layers.

Layer	Components	Role
1. Hardware	IBM Power11 (Power S1112 entry node; larger Power11 scale-up nodes for departmental/enterprise tiers); optional IBM Spyre PCIe AI accelerator; IBM LinuxONE option for extreme-resiliency banking/telco workloads; NVMe/Red Hat Ceph storage; resilient power and connectivity kit for county/edge sites	On-chip Matrix Math Acceleration for local inference without a mandatory discrete GPU; Spyre adds scalable, low-power off-chip acceleration for heavier vertical workloads
2. Platform	Red Hat OpenShift AI (RHOAI) on Red Hat Enterprise Linux; vLLM inference serving; Red Hat AI model-serving runtime	Kubernetes-native, hybrid-portable base that lets the same appliance image run on-prem, collocated or in IBM Cloud for burst capacity
3. Orchestration & governance	watsonx Orchestrate (multi-agent control plane) + IBM Sovereign Core (runtime policy and compliance enforcement) + IBM Concert (AI-powered operations) + IBM Bob (agentic development assistant for extending the appliance)	Deploys and governs agents from any vendor under one policy layer; embeds compliance controls at the infrastructure level rather than as an application-level bolt-on
4. Model garden	Open-source general-reasoning LLMs (Llama/Mistral-class) for broad-domain tasks; Granite 4.1/4.2 family — language (3B/8B/30B), vision (document extraction), speech (transcription), embedding (retrieval) and Guardian (safety) models, all Apache-2.0 and cryptographically signed	General-purpose reasoning paired with small, auditable, cost-efficient specialists for the narrow tasks that make up most enterprise workload volume
5. Vertical intelligence	i3's own fine-tuned specialty products — AgriGranite, FinGranite, GovGranite and the balance of the 16-product Vertical AI Solutions Blueprint — plus Kenyan-language and Swahili/Sheng speech models from the Royal Partner Program / IBM Research Lab Africa collaboration	The commercial differentiator: the layer competitors without i3's existing blueprint and partner network cannot easily replicate

Because every layer above the hardware is already portable across Red Hat OpenShift, the same certified image can be shipped as a physical appliance, deployed inside an i3/partner colocation facility, or installed directly inside a customer's own data centre — with IBM Cloud available as an optional burst layer for non-sensitive, elastic workloads. This hybrid portability is the same characteristic IBM highlighted at Think 2026 for regulated buyers seeking sovereign-data handling compatible with grant programmes such as India's IndiaAI Mission — the identical pattern applies to Kenyan and East African sovereign-AI funding and county PPP structures.

Important correction: Sovereign Core is hardware-agnostic — design i3-SAA accordingly

IBM Sovereign Core, as it actually reached general availability in May 2026, is explicitly designed to run “on customer-provided infrastructure across compute, storage and network layers” — it is a governance and control-plane layer, not a Power-exclusive stack. IBM's own published Sovereign Core ecosystem currently lists AMD, ATOS, Cegeka, Cloudera, Dell, Elastic, HCL, Intel, Mistral, MongoDB and Palo Alto Networks as launch partners — AMD is a named compute launch partner alongside IBM Power, and the catalogue is built to be

extensible with pre-vetted third-party and open-source components.

The appliance should therefore be positioned as **“Sovereign Core certified, IBM Power-optimised” rather than IBM Power-exclusive**. This has two practical benefits: it widens i3's addressable market to customers or IBM engagements that specify AMD- or Dell-based compute, and it aligns i3's own certification path directly with the IBM Partner Plus / Sovereign Core ecosystem programme rather than with a narrower Power-hardware resale relationship. Power11/Spyre remains the recommended and reference configuration — it is where i3's existing ROKS/RHOAI AI Lab experience and IBM Silver Partner relationship already sit — but the appliance's software and vertical-model layers should be validated to run on any Sovereign Core–certified infrastructure.

5. Product Line: Appliance Tiers (SKUs)

Following the appliance-market discipline of a small number of fixed, well-understood SKUs rather than unlimited custom configurations, i3 should launch with four tiers. Each tier is quoted below on its reference IBM Power11 configuration; because the software stack is Sovereign Core–certified rather than Power-exclusive (see Section 4), each tier can also be fulfilled on a certified AMD- or Dell-based configuration where a customer's existing estate or an IBM-led engagement specifies it.

Tier	Target buyer	Hardware profile	Typical workload
i3-SAA Edge	Branch office, county sub-office, single clinic or SACCO branch	1× Power S1112 entry node, on-chip MMA only	Single-language citizen/customer desk, document intake, one or two vertical SLMs
i3-SAA Departmental	Single department of a bank/insurer/hospital, or a county headquarters	1–2 Power11 nodes + Spyre accelerator	Multi-agent customer operations, RAG knowledge assistant, moderate concurrent users
i3-SAA Enterprise	Full enterprise deployment — banks, insurers, telcos, national ministries	Clustered Power11 nodes, multiple Spyre accelerators, high-availability storage; LinuxONE option	Full AI-BPO tower, multiple concurrent vertical products, 24/7 SLA-bound operations
i3-SAA Sovereign Cloud-in-a-Box	Government agencies, regulated national infrastructure	Redundant multi-node Power11/LinuxONE cluster with full Sovereign Core policy stack	Multi-tenant platform serving several agencies/counties from one governed, sovereign environment

Each tier ships with the same certified software and governance stack — only the hardware footprint and concurrency envelope change — so a customer that starts on Edge can grow into Departmental or Enterprise without a re-platforming project. This is the same “land small, expand in place” motion that made appliance vendors durable annuity businesses rather than one-off hardware sellers.

6. Vertical Industry Model Library

The appliance's real moat is not the hardware — any competent systems integrator can rack a Power11 server. It is the pre-loaded, pre-evaluated vertical intelligence layer i3 has already begun building through its Granite-based Vertical AI Solutions Blueprint (AgriGranite, FinGranite, GovGranite and 13 further named products across 16 total verticals, mapped to Red Hat OpenShift infrastructure and IBM credit bands). The appliance programme should treat that blueprint as its factory-installed software catalogue.

Vertical pack	Representative capabilities on the appliance
FinGranite — Banking, SACCOs, Microfinance	KYC/AML case triage, loan-document intake, credit-file summarisation, collections prioritisation, fraud/anomaly investigation support — all processed on-premises inside the bank's own risk boundary
Insurance pack	Claims intake and classification, policy Q&A copilots, renewal/lapse prevention, fraud red-flag queues
GovGranite — County & national government	Citizen information desks, document intake, public-participation summarisation, county knowledge assistants trained only on approved public records
AgriGranite — Agriculture & agribusiness	Farmer voice advisory in Swahili/Sheng/local languages, cooperative call operations, market-price dissemination, crop-record digitisation
Healthcare pack	Patient navigation, medical-document summarisation with clinician oversight, insurance pre-authorisation support
Manufacturing / Energy pack	Maintenance copilots, SOP/safety knowledge assistants, work-order classification, technician training simulations
Education pack	AI tutors, assessment generation and moderation, institution knowledge assistants — aligned to the existing i3 Academy training catalogue
Language & Voice pack	Kenyan-English, Swahili and Sheng speech-to-text/text-to-speech built from the Royal Partner Program and IBM Research Lab Africa collaboration at CUEA

Each pack runs on Granite's vision, speech, embedding and Guardian safety models alongside the core language models — the same combination IBM ships as one signed, auditable family — so a single appliance can, for example, extract a scanned loan document (vision), summarise it (language), and check the output against safety and compliance rules (Guardian) without any component leaving the customer's network.

7. Three Delivery Motions

The proposal brief specifically asks for delivery “to customer, or collocated, or in their data centre.” All three should be first-class, catalogue-priced options from launch.

Motion	Description	Best fit	i3's role
Ship-to-site appliance	Factory-configured appliance shipped and installed at the customer's own premises	Branches, clinics, SACCOs, county sub-offices — lower initial data-centre maturity	Manufacture/configure, ship, install, remote-manage under SLA
i3/partner colocation	Appliance racked at an i3-controlled or partner colocation facility (e.g. Nairobi data-centre capacity such as iXAfrica-class facilities) with dedicated, logically isolated tenancy	Customers without their own data-centre space, or wanting managed physical security without losing sovereignty	Own and operate the facility; sell dedicated capacity under a hosting + managed-service contract
Customer data-centre deployment	i3 installs and operates the appliance inside the customer's existing data centre, fully under the customer's physical control	Banks, insurers, government ministries with strict data-residency and physical-custody requirements	Install, integrate, and provide remote/on-site managed operations under SLA; customer retains physical custody

All three motions share the same certified software image and the same watsonx Orchestrate/Sovereign Core governance layer, so a customer's data-residency posture never has to trade off against the sophistication of the AI stack — they get the identical vertical intelligence library regardless of where the physical box sits.

8. Managed-Service Ladder (Appliance-as-a-Service)

Consistent with i3's existing packaging logic for its broader AI Services Utility, the appliance should never be sold as a one-off hardware transaction. A managed-service ladder converts each sale into a recurring annuity and gives i3 the operational data (lawfully governed) needed to keep improving the vertical model library.

Package	What it includes	Primary KPI
Sovereignty Readiness Assessment	Data-residency and workload audit; appliance-tier sizing; ROI case	Validated appliance sizing & business case
Appliance Pilot	One appliance, one or two vertical packs, in production for 60–90 days	Task success rate, cost/time saved
Managed Appliance Operations	24/7 monitoring, patching, model updates, human-in-the-loop QA, IBM Power Autonomous Operations integration	SLA uptime, evaluation pass rate
Vertical Pack Subscription	Ongoing access to updated AgriGranite/FinGranite/GovGranite and related packs as they improve	Model quality/version currency
AI Evaluation-as-a-Service	Continuous benchmarking, hallucination/safety/regression testing via Guardian models	Evaluation pass rate, incident rate
Sovereign Cloud-in-a-Box Programme	Multi-agency/county shared appliance with tenancy isolation and unified governance	Agencies onboarded, resolution rate

9. Regulatory & Governance Alignment

The appliance model is designed around, not around the outside of, Kenya's and the region's emerging AI and data-protection regime — which is itself the appliance's strongest sales argument.

Requirement	How the appliance addresses it
Kenya Data Protection Act, 2019 / ODPC controller-processor registration	Appliance ships with a data-provenance register and consent-ledger pattern; i3 supports customers through ODPC controller/processor registration as part of onboarding
Kenya cloud policy — government data stored within national territory	On-premises and in-country colocation motions satisfy residency by construction — data never crosses a border unless the customer explicitly configures a burst path
Kenya National AI Strategy 2025–2030 / draft AI Bill (proposed AI Commissioner, high-risk audits, 5-year retention)	Sovereign Core's runtime policy enforcement and Guardian safety models generate the audit trail a future AI Commissioner would expect to see
Sectoral data-localisation for health and financial records	Vertical packs for banking, insurance and healthcare are designed to run entirely on the on-prem/colocated appliance, never routing sensitive fields to external inference
African Union Continental AI Strategy — data governance and sovereignty pillar	The appliance keeps agricultural, health and citizen data generated in-country processed in-country, directly answering the AU's stated concern
EU AI Act Article 10 (for exporters/partners with EU exposure)	Model cards, Guardian evaluation and Sovereign Core audit logging give exportable evidence of data-governance documentation where relevant

As with the broader AI Services Utility strategy, the governing principle remains evidence, provenance and human escalation: AI confidence is not evidence, and every appliance deployment should distinguish verified, partially verified, unverified and contradicted outputs, with human sign-off retained for legal, medical, financial, identity and other high-impact decisions.

10. Competitive Positioning & Moat

Three categories of competitor will contest this space, and the appliance model is deliberately structured to be defensible against each.

Competitor type	Their likely offer	i3's counter-position
Global hyperscalers (AWS/Azure/Google Cloud AI)	Powerful managed AI services, but cloud-resident by default and subject to foreign jurisdiction and CLOUD-Act-style exposure	i3 sells data that never leaves the customer's building or country — the exact gap hyperscaler regional pushes cannot close as fast as an on-prem appliance can
Generic systems integrators / hardware resellers	Can rack IBM Power hardware, but have no vertical model library and no evaluation/governance layer	i3 owns the Granite-based vertical blueprint (AgriGranite/FinGranite/GovGranite and the rest of the 16-product catalogue) that a reseller cannot replicate quickly
International BPO/AI-ops vendors	Offer AI-enabled outsourcing, but typically from offshore delivery centres, not sovereign in-country infrastructure	i3 combines the appliance with local workforce (AI Operator Academy, Royal Partner Program) to deliver both sovereign infrastructure and local human-in-the-loop capacity
Regional "AI Factory" builders (Cassava Technologies / NVIDIA Cloud Partners)	Hyperscale, multi-tenant NVIDIA GPU capacity sold as regional cloud, with an expansion roadmap that explicitly names Kenya	i3 does not compete on raw GPU capacity; it sells governed, appliance-sized, single-tenant units for the vertical, regulated workloads a generic GPU cloud does not serve — and can treat their in-country facilities as a colocation option, not only a rival

The durable moat is the same one identified in i3's broader AI Services Utility strategy: not any single model or hardware vendor, but the combined stack of people, proprietary evaluations, local language capability, existing customer relationships (via the Royal Partner Program's 65 nodes across 11+ countries), and a certified, pre-integrated product that removes the buyer's integration risk.

10.1 Africa's "AI Factory" race: a validated category, and where i3-SAA fits inside it

The sovereign-AI-appliance thesis is no longer a novel idea i3 would be pioneering alone — it is now a named global category, with NVIDIA running the same "national or regional champion plus hyperscale compute plus local model layer" playbook across multiple continents in 2026: NAVER with NVIDIA's DSX platform in South Korea, TELUS in Canada, Deutsche Telekom's Industrial AI Cloud in Germany, and SUSE's turnkey AI Factory software stack in Europe. That global validation is useful — boards and regulators increasingly recognise the model — but it has already arrived in East Africa under a different banner.

Cassava Technologies (part of Strive Masiyiwa's Econet group) became Africa's first NVIDIA Cloud Partner and brought an NVIDIA-powered AI factory online in South Africa in March 2026, with a publicly stated expansion roadmap that explicitly names **Kenya**, alongside Nigeria, Egypt and Morocco; a further factory with Vodafone Egypt went live in Cairo in September 2026. This is real, funded, government-adjacent infrastructure, not a competitor to be dismissed.

The right response is not to out-build Cassava's compute scale — that is a capital-intensive, GPU-hyperscale game i3 is not positioned to win. It is to occupy the adjacent, complementary layer their model does not serve well: single-tenant, governed, vertical-industry appliances sized for one bank branch, one county office, one hospital — built on IBM's Sovereign Core/Power stack rather than a shared NVIDIA GPU pool, and pre-loaded with the AgriGranite/FinGranite/GovGranite vertical intelligence layer a generic compute-cloud provider has no reason to

build. If and when Cassava's Nairobi capacity lands, it is worth a direct conversation about i3-SAA units being colocated inside their facility — turning a potential rival's infrastructure investment into i3's colocation delivery motion (Section 7).

11. Go-to-Market: Kenya -> East Africa -> MEA

The rollout should prove the appliance model in a small number of reference accounts in Kenya before using the existing partner network to replicate regionally — the same discipline recommended for the broader Digital Work-Center programme: prove revenue in a repeatable unit before scaling infrastructure.

Phase	Geography	Focus
Phase 1 (0–6 months)	Kenya — Nairobi	Build the reference i3-SAA Departmental appliance on the existing ROKS/RHOAI AI Lab; certify 2–3 vertical packs (FinGranite, GovGranite); land 2–3 paid pilots with a bank/SACCO, a county government and one insurer
Phase 2 (6–12 months)	Kenya — national	Convert pilots to managed-service contracts; launch the Edge tier for county sub-offices and SACCO branches; formalise the colocation motion at a Nairobi facility
Phase 3 (12–24 months)	East Africa (Uganda, Tanzania, Rwanda, South Sudan, Somalia, Burundi)	Deploy through the Royal Partner Program's regional nodes; adapt vertical packs to local language and regulatory variants; add Rwanda's National AI Agency and similar counterparts as government channels
Phase 4 (24–36 months)	Middle East & Africa (MEA)	Use IBM's regional alliance structure and the Anthropic partnership to enter Gulf and North/West African sovereign-AI programmes, where sovereign-cloud demand (an estimated USD 80B global market, growing 35.6% y/y) is already well established

12. Approaching IBM: Two Separate Tracks, Not One Conversation

IBM's Sovereign Core organisation and IBM Research Africa are different parts of IBM, with different objectives, different currencies of value, and different points of contact. i3 should run both relationships in parallel rather than routing everything through a single account manager, and should not expect one relationship to open both doors.

Track A — Commercial: IBM Partner Plus / Sovereign Core ecosystem

IBM has been explicit that Sovereign Core “gives business partners a way to build repeatable deployments in multi-tenant models they operate” — which is, almost verbatim, the i3-SAA product thesis, offered as an official IBM programme rather than something i3 would have to build unsanctioned. IBM's published Sovereign Core ecosystem currently lists AMD, ATOS, Cegeka, Cloudera, Dell, Elastic, HCL, Intel, Mistral, MongoDB and Palo Alto Networks as partners — **none headquartered in or focused on Africa**. That is a specific, nameable gap i3 can ask to fill.

- ◆ Use the existing IBM Silver Business Partner relationship to request a direct introduction to the Sovereign Core / Partner Plus ecosystem owner covering Africa — the ask should be narrow and specific: “we want to become IBM's first African Sovereign Core ecosystem partner,” not a general partnership renewal conversation.
- ◆ Bring this proposal's appliance tiers, delivery motions and vertical-pack catalogue as the business case — IBM's own materials show they are actively building out the Sovereign Core catalogue with “pre-vetted IBM, third-party and open-source software and services,” which is exactly the slot AgriGranite/FinGranite/GovGranite can fill.
- ◆ In parallel, pursue Red Hat Authorised Training Partner status as the credentialing path for an AI Operator Academy that produces IBM/Red Hat–certified delivery staff — a concrete asset IBM's ecosystem team can point to when evaluating i3 as a regional partner.

Track B — Research: IBM Research Africa, Nairobi

IBM Research Africa is a distinct organisation that trades in research value, not licence revenue — its track record in Kenya includes a formal Memorandum of Understanding with the Office of the Special Climate Envoy applying watsonx.ai geospatial models to water-tower restoration, alongside an active internship and entrepreneur-engagement programme. The way in is a specific, fundable research proposal, not a general partnership pitch.

- ◆ Propose a joint initiative on **Swahili and Sheng specialty small language models** — directly reusable in the appliance's Language & Voice pack — offering i3's engineering bench (75 AI/ML engineers today, ramping to 150 by March 2027) as embedded co-development capacity working alongside IBM Research Africa's team.
- ◆ Frame it as a trade, not a request: i3 brings African-language data, distribution and delivery hands; IBM Research brings Granite model research depth, evaluation tooling and credibility — the same shape as the existing Kenya climate MoU, applied to language technology instead of geospatial data.
- ◆ Keep this proposal separate from the Track A commercial conversation — a research MoU with IBM Research Africa does not require, and should not wait for, progress on the Sovereign Core ecosystem application, and vice versa.

Existing i3 assets that strengthen both tracks

- ◆ **The Royal Partner Program** (65 nodes across 11+ African countries, expanding to South Africa, Somalia, South Sudan and Burundi) becomes the regional distribution and installation network for the appliance once the Kenyan reference deployment is proven — a concrete regional footprint IBM's ecosystem team can evaluate.
- ◆ **The existing collaboration framework with IBM Research Lab Africa at CUEA** is the natural home for the Track B Swahili/Sheng SLM proposal, rather than a cold approach to the Nairobi lab.

- ◆ **Anthropic Partner status** extends the appliance's model garden beyond IBM Granite to Claude models where a customer's use case calls for frontier-class reasoning alongside on-device Granite specialists — subject to the relevant commercial and deployment terms.
- ◆ **The existing Vertical AI Solutions Blueprint and its ~92-organisation outreach tracker** already identify a warm pipeline of Kenyan target accounts that can be re-approached with the appliance as a concrete, shippable product rather than a conceptual roadmap.

13. Top Risks & Mitigations

Risk	Severity	Mitigation
Watsonx Orchestrate next-gen and Sovereign Core partner ecosystem remain in preview/early GA through 2026	Medium	Sequence the roadmap so Phase 1 uses currently GA components (Power11, Sovereign Core GA, Granite 4.1) and treats next-gen Orchestrate features as a Phase 2 upgrade, not a launch dependency
Hardware supply/lead time for Power11 and Spyre accelerators into East Africa	Medium	Pre-order a small reference fleet for the Kenyan pilots; use the IBM Silver Partner relationship to secure allocation priority
Data-governance missteps with sensitive sectoral data (health, financial, identity)	Critical	Data Protection Impact Assessments, controller/processor registration with ODPC, and Sovereign Core policy enforcement built into every deployment from day one
Regulatory uncertainty while Kenya's AI Bill and Implementation Secretariat are still forming	Medium	Design to the strictest plausible reading of the draft Bill now, so later formalisation is a compliance confirmation rather than a rebuild
Commoditisation — competitors reselling bare Power hardware without the vertical layer	Medium	Protect the vertical model library and evaluation IP as the differentiator; do not sell hardware without the managed-service and model layer attached
Talent constraint for Power11/OpenShift AI specialists in-market	Medium	Use Red Hat ATP status and i3 Academy to certify an internal and partner-network talent pipeline ahead of scale-out

14. First 12 Months: Build -> Prove -> Scale

Window	Milestone
Month 0–2	Stand up the reference i3-SAA Departmental appliance on the existing ROKS/RHOAI AI Lab; formalise IBM Sovereign Core and watsonx Orchestrate access via the Silver Partner relationship
Month 2–4	Certify FinGranite and GovGranite packs on the reference appliance; complete Data Protection Impact Assessments for the first two pilot accounts
Month 4–6	Launch two paid pilots — one financial-sector, one county/government — using the ship-to-site and colocation motions respectively
Month 6–9	Convert pilots to managed-service contracts; launch AI Operator Academy cohort; begin Edge-tier certification for branch/sub-office deployment
Month 9–12	Publish reference case studies; open the Royal Partner Program channel for the first East African replication (Uganda or Tanzania); begin MEA partner conversations via IBM's regional alliance structure

15. Recommended Positioning

i3 Technologies should position itself as **Africa's Sovereign AI Appliance company** — the vendor that turns IBM's 2026 sovereign-AI hardware and software roadmap, combined with i3's own Granite-based vertical intelligence and pan-African partner network, into a product a Kenyan bank, county government or hospital can order, receive, and run entirely within its own walls. The appliance line should sit as a named product family inside i3's broader AI Services Utility strategy — the physical, ownable counterpart to the managed-services and BPO layers already under development — with its own roadmap, SKUs and P&L discipline, launched first in Kenya, proven through the existing partner and outreach network across East Africa, and extended into the Middle East & Africa region once the model is proven.

Final recommendation: commission a focused engineering sprint to certify the i3-SAA Departmental reference appliance on currently generally-available IBM components (Power11, Sovereign Core, Granite 4.1) within the next 60 days, in parallel with formal Data Protection Impact Assessments for the first two target pilot accounts — so that the appliance's first paid, production deployment lands inside Phase 1 of the roadmap above.

This document is a strategic and technical proposal, not a binding commercial quotation, a regulatory determination, or a confirmation that any specific IBM programme, partner term, or model licence is available to i3 on the terms assumed here. Hardware, software and partner-programme specifics should be confirmed directly with IBM, Red Hat and Anthropic account teams before customer commitments are made.